

Neural Networks in OpenCL

Jakub Zárybnický

Implementation

- My own subgoal: result in Haskell
- CUDA
 - nice API, pleasant bindings but not emulable
 - lost access to Gent datacenter @Dec 10 due to renovations
- Raw OpenCL first
- OpenCL in Haskell (but then got uncomfortably close to the existing Grenade API)
- Parallelizing an existing framework (Grenade)
 - Adding an OpenCL backend

Parallel v. original gradient descent

```
kernel void descend_slow(  
    double rate,  
    double momentum,  
    double regulariser,  
    constant double *weights,  
    constant double *gradient,  
    constant double *last,  
    global double *outputWeights,  
    global double *outputMomentum  
) {  
    int i = get_global_id(0);  
    double momentum_ = momentum * last[i] - rate * gradient[i];  
    outputMomentum[i] = momentum_;  
    outputWeights[i] = mad(weights[i], 1 - rate * regulariser, momentum_);  
}  
  
-:--- kernel.cl      Top (9,33)      Git-master (Opencl/*l FlyC Projectile[gmu-project:haske  
#include "gradient_descent.h"  
  
void descend_cpu(int len, double rate, double momentum, double regulariser,  
    const double* weights,  
    const double* gradient,  
    const double* last,  
    double* outputWeights, double* outputMomentum) {  
  
    for (int i = 0; i < len; i++) {  
        outputMomentum[i] = momentum * last[i] - rate * gradient[i];  
        outputWeights[i] = weights[i] + outputMomentum[i] - (rate * regulariser) * weights[i];  
    }  
}
```

FFI API replacements

```
clEnqueueWriteBuffer' cq mem check off size dat =
  alloca $ \event → do
    _ ← raw_clEnqueueWriteBuffer cq mem (fromBool check) (fromIntegral off) (fromIntegral size)
    raw_clWaitForEvents 1 (castPtr event)

clEnqueueNDRangeKernel' :: CLCommandQueue → CLKernel → [CSize] → IO Int32
clEnqueueNDRangeKernel' cq krn gws =
  withArray gws $ \pgws →
    alloca $ \event →
      whenSuccess
        (raw_clEnqueueNDRangeKernel cq krn num nullPtr pgws nullPtr 0 nullPtr event)
        (raw_clWaitForEvents 1 (castPtr event))
  where
    num = fromIntegral $ length gws

clSetKernelArgSto' :: Storable a ⇒ CLKernel → Word32 → a → IO ()
clSetKernelArgSto' krn idx val = with val $ \pval → do
  () <$ raw_clSetKernelArg krn idx (fromIntegral . sizeof $ val) (castPtr pval)

foreign import ccall "clEnqueueNDRangeKernel" raw_clEnqueueNDRangeKernel
  :: CLCommandQueue
  → CLKernel
  → Word32
  → Ptr CSize
  → Ptr CSize
  → Ptr CSize
  → Word32
  → Ptr CLEvent
  → Ptr CLEvent
  → IO Int32

foreign import ccall "clEnqueueReadBuffer" raw_clEnqueueReadBuffer
  :: CLCommandQueue
  → CLMem
  → CBool
  → CSize
  → CSize
  → Ptr ()
  → Word32
  → Ptr CLEvent
  → Ptr CLEvent
```

Technologies

- Haskell
- Nix
- OpenCL bindings
- Grenade NN framework
 - Internally uses Hmatrix (= BLAS/LAPACK) for matrix operations
- (formerly also CUDA, raw OpenCL, ...)

Desired interface

```
type MNIST
  = Network
    '[ Convolution 1 10 5 5 1 1, Pooling 2 2 2 2, Relu
      , Convolution 10 16 5 5 1 1, Pooling 2 2 2 2, Reshape, Relu
      , FullyConnected 256 80, Logit, FullyConnected 80 10, Logit]
    '[ 'D2 28 28, 'D3 24 24 10, 'D3 12 12 10, 'D3 12 12 10
      , 'D3 8 8 16, 'D3 4 4 16, 'D1 256, 'D1 256
      , 'D1 80, 'D1 80, 'D1 10, 'D1 10]

randomMnist :: MonadRandom m => m MNIST
randomMnist = randomNetwork
```

Network building block

```
data Network :: [*] -> [Shape] -> * where
  NNil  :: SingI i
        => Network '[] '[i]

  (:~>) :: (SingI i, SingI h, Layer x i h)
        => !x
        -> !(Network xs (h ': hs))
        -> Network (x ': xs) (i ': h ': hs)
```

Test network

```
type FFNet = Network
  '[ FullyConnected 2 40
    , Tanh
    , FullyConnected 40 10
    , Relu
    , FullyConnected 10 1
    , Logit
  ]
  '[ 'D1 2
    , 'D1 40
    , 'D1 40
    , 'D1 10
    , 'D1 10
    , 'D1 1
    , 'D1 1
  ]
```


Test network API

```

. .-==###==.
-#####=.
.#####=.
.#####=.
-#####=.
-#####=-.
. -#####=-.
. --=##=-. . .--.
. -#####=-.
.=#####-.
=#####-
.#####=.
.#####-
. -#####=.

```

Training API

```
randomNet :: (MonadIO m, MonadRandom m) => m FFNet
randomNet = randomNetwork

netTrain :: FFNet -> LearningParameters -> Int -> IO FFNet
netTrain net0 rate n = do
  inps <- replicateM n $ do
    s <- getRandom
    return $ S1D $ SA.randomVector s SA.Uniform * 2 - 1
  let outs = flip map inps $ \(S1D v) ->
    if v `inCircle` (fromRational 0.33, 0.33) || v `inCircle` (fromRational (-0.33), 0.33)
    then S1D $ fromRational 1
    else S1D $ fromRational 0

  let trained = foldl' trainEach net0 (zip inps outs)
  return trained
```

Accelerated layer internals

```
data FullyConnectedCL i o = FullyConnectedCL
  { kUpdate :: !CLKernel
  , kForward :: !CLKernel
  , kBackward :: !CLKernel
  , lpRef :: !(IORef (Maybe LearningParameters))
  , wTape :: !(CLBuffer (R i))          -- work vector
  , wIn :: !(CLBuffer (R i))            -- work vector
  , wOut :: !(CLBuffer (R o))           -- work vector
  , wGradient :: !(FullyConnected'CL i o) -- work gradient
  , wWeights :: !(FullyConnected'CL i o) -- Neuron weights
  , wMomentum :: !(FullyConnected'CL i o) -- Neuron momentum
  }

data FullyConnected'CL i o = FullyConnected'CL
  !(CLBuffer (R o)) -- Bias
  !(CLBuffer (L o i)) -- Activations
```

```

randomFullyConnectedCL :: forall i o m. (MonadIO m, MonadRandom m, KnownNat i, KnownNat o)
    => m (FullyConnectedCL i o)
randomFullyConnectedCL = do
  s1 <- getRandom
  s2 <- getRandom
  liftIO . withCL $ \cl -> do
    oB <- mkBufferR cl (randomVector s1 Uniform * 2 - 1)
    oA <- mkBufferL cl (uniformSample s2 (-1) 1)
    oBM <- mkBufferR cl (konst 0)
    oM <- mkBufferL cl (konst 0)
    bG <- mkBufferR cl (konst 0)
    aG <- mkBufferL cl (konst 0)
    wIn <- mkBufferR cl (konst 0)
    wOut <- mkBufferR cl (konst 0)
    wTape <- mkBufferR cl (konst 0)
    let wWeights = FullyConnected'CL oB oA
        wMomentum = FullyConnected'CL oBM oM
        wGradient = FullyConnected'CL bG aG

    kUpdate <- clCreateKernel (clProgram cl) "fcnn_update"
    clSetKernelArgSto' kUpdate 0 (fromIntegral @_ @Int32 (natVal (Proxy @o)))
    clSetKernelArgSto' kUpdate 1 (fromIntegral @_ @Int32 (natVal (Proxy @i)))
    clSetKernelArgSto' kUpdate 5 oB
    clSetKernelArgSto' kUpdate 6 oA
    clSetKernelArgSto' kUpdate 7 bG
    clSetKernelArgSto' kUpdate 8 aG
    clSetKernelArgSto' kUpdate 9 oBM
    clSetKernelArgSto' kUpdate 10 oM

    kForward <- clCreateKernel (clProgram cl) "fcnn_forward"
    clSetKernelArgSto' kForward 0 (fromIntegral @_ @Int32 (natVal (Proxy @o)))
    clSetKernelArgSto' kForward 1 (fromIntegral @_ @Int32 (natVal (Proxy @i)))
    clSetKernelArgSto' kForward 2 oA
    clSetKernelArgSto' kForward 3 wTape
    clSetKernelArgSto' kForward 4 oB
    clSetKernelArgSto' kForward 5 wOut

    kBackward <- clCreateKernel (clProgram cl) "fcnn_backward"
    clSetKernelArgSto' kBackward 0 (fromIntegral @_ @Int32 (natVal (Proxy @o)))
    clSetKernelArgSto' kBackward 1 (fromIntegral @_ @Int32 (natVal (Proxy @i)))
    clSetKernelArgSto' kBackward 2 bG
    clSetKernelArgSto' kBackward 3 wTape
    clSetKernelArgSto' kBackward 4 oA
    clSetKernelArgSto' kBackward 5 aG
    clSetKernelArgSto' kBackward 6 wIn

  lpRef <- newIORef Nothing

  pure $ FullyConnectedCL{..}

```

Backpropagation kernel

```
kernel void fcnn_backward(  
    int rows,  
    int cols,  
    constant double *dEdy,  
    constant double *x,  
    constant double *wN,  
    global double *mm,  
    global double *dWs  
) {  
    int i = get_global_id(0);  
  
    // let mm' = dEdy `outer` x  
    if (i < rows) {  
        const double d = dEdy[i];  
        for (unsigned j = 0; j < cols; j++) {  
            mm[i * rows + j] = d * x[j];  
        }  
    }  
  
    // dWs = tr wN #> dEdy  
    if (i < cols) {  
        double acc = 0;  
        for (unsigned j = 0; j < rows; j++) {  
            acc = mad(wN[i + rows * j], dEdy[j], acc);  
        }  
        dWs[i] = acc;  
    }  
}
```

Bottleneck – readVectorFromBuffer

```
readBufferL :: forall m n. (KnownNat m, KnownNat n) => CLState -> CLBuffer (L m n) -> Maybe (L m n) -> IO (L m n)
readBufferL cl (CLBuffer buf) mMat = do
  fPtr <- case mMat of
    Nothing -> mallocForeignPtrArray len
    Just mat -> pure $ fst (unsafeToForeignPtr0 (flatten . tr $ extract mat))
  _ <- withForeignPtr fPtr $ \ptr ->
    clEnqueueReadBuffer' (clQueue cl) buf True 0 size (castPtr ptr)
  pure (fromJust . create . matrixFromVector ColumnMajor rows cols $ unsafeFromForeignPtr0 fPtr len)
  where
    rows = fromIntegral (natVal (Proxy @m))
    cols = fromIntegral (natVal (Proxy @n))
    len = rows * cols
    size = sizeOf (undefined :: CDouble) * len
```



```

unsafeWithCL $ cl → do
    individual
    inherited
let k = noIte entries %time %alloc %time %alloc
lastIP ← readIORef (ipRef n)
unless (317 ip == lastIP) do
    clSetKey 633 ArgSto k 0 0.0 0.0 0.0 100.0 100.0
    clSetKey 674 ArgSto k 1 0.0 0.0 0.0 0.0 0.0
    clSetKey 675 ArgSto k 1 0.0 0.0 0.0 0.0 0.0
writeIORef (ipRef n) 676 0.0 0.0 0.0 0.0 0.0
:1-28 -- clEnqueueRange 680 4 0.0 0.0 0.0 0.0
(65) -- clEnqueueRange 681 4 0.0 0.0 0.0 0.0
(3,1)-(167,36) 681 4 0.0 0.0 0.0 0.0
len = fromIntegral (max 4 1) 679 0.0 0.0 0.0 0.0
677 2 0.0 0.0 0.0 0.0
:1-28 clEnqueueRange 678 2 0.0 0.0 0.0 0.0
682 2 0.0 0.0 0.0 0.0
(65) clEnqueueRange 683 1 0.0 0.0 0.0 0.0
(3,1)-(167,36) 684 2 0.0 0.0 0.0 0.0
634 1 0.0 0.0 0.0 0.0
-- do a matrix 647 cl ← do
(34) (v, SIO w) 646 1 0.0 0.0 0.0 0.0
runForwards n 808 0 v) 0 0.0 0.0 0.0 0.0
unsafeWithCL 818 cl → do
    writeBuff 822 cl (wType 7 v) 0.0 0.0 0.0 0.0
    -- clEnqueueRange 821 2 0.0 0.0 0.0 0.0
res = readIORef cl 820 2 0.0 0.0 0.0 0.0
(30) pure (wType 819, SIO res) 2 0.0 0.0 0.0 0.0
where
    len = fromIntegral (max 4 1) 632 0.0 0.0 0.0 0.0
(54) 692 1 0.0 0.0 0.0 0.0
:1-25 writeIORef 693 0 0.0 0.0 0.0 0.0
(16) let wType 697 0 0.0 0.0 0.0 0.0
-- do a matrix 699 wType 1 0.0 0.0 0.0 0.0
:1-25 n = (fully 700 wType) 0 0.0 0.0 0.0 0.0
runBackwards 698 (SIO dEdge) 1 0.0 0.0 0.0 0.0
(65) unsafeWithCL 663 cl → do
    let fully 664 readIORef cl 1 0.0 0.0 0.0 0.0
:1-25 writeBuff 665 cl (dEdge 0 0.0 0.0 0.0 0.0
    -- clEnqueueRangeKernel (clQueue cl) (kBackward n) [len]
dEdge = readIORef cl (wType n) (Just (konst 0))

```


runNet	Grenade.Core.Runner	src/Grenade/Core/Runner.hs:57:1-33	809	1071	0.0	0.0	0.7	0.8
runNetwork	Grenade.Core.Network	src/Grenade/Core/Network.hs:(103,1)-(116,17)	811	0	0.4	0.5	0.7	0.8
unId	Data.Vector.Fusion.Util	Data/Vector/Fusion/Util.hs:25:21-24	867	34272	0.0	0.0	0.0	0.0
clEnqueueNDRangeKernel'	Grenade.OpenCL.Context	src-gmu/Grenade/OpenCL/Context.hs:(229,1)-(236,35)	859	3213	0.0	0.0	0.0	0.0
clQueue	Grenade.OpenCL.Context	src-gmu/Grenade/OpenCL/Context.hs:59:5-11	860	3213	0.0	0.0	0.0	0.0
kForward	Grenade.OpenCL.FullyConnected	src-gmu/Grenade/OpenCL/FullyConnected.hs:29:5-12	861	3213	0.0	0.0	0.0	0.0
readBufferR	Grenade.OpenCL.Context	src-gmu/Grenade/OpenCL/Context.hs:(132,1)-(141,46)	862	3213	0.2	0.1	0.2	0.2
unId	Data.Vector.Fusion.Util	Data/Vector/Fusion/Util.hs:25:21-24	863	12852	0.0	0.0	0.0	0.0
clEnqueueReadBuffer'	Grenade.OpenCL.Context	src-gmu/Grenade/OpenCL/Context.hs:(216,1)-(219,41)	866	3213	0.0	0.0	0.0	0.0
clQueue	Grenade.OpenCL.Context	src-gmu/Grenade/OpenCL/Context.hs:59:5-11	864	3213	0.0	0.0	0.0	0.0
sChunks	Data.Vector.Fusion.Bundle.Monadic	Data/Vector/Fusion/Bundle/Monadic.hs:122:30-36	870	3213	0.0	0.0	0.0	0.0
sSize	Data.Vector.Fusion.Bundle.Monadic	Data/Vector/Fusion/Bundle/Monadic.hs:124:30-34	869	3213	0.0	0.0	0.0	0.0
upperBound	Data.Vector.Fusion.Bundle.Size	Data/Vector/Fusion/Bundle/Size.hs:(126,1)-(128,30)	868	3213	0.0	0.0	0.0	0.0
wOut	Grenade.OpenCL.FullyConnected	src-gmu/Grenade/OpenCL/FullyConnected.hs:34:5-8	865	3213	0.0	0.0	0.0	0.0
wTape	Grenade.OpenCL.FullyConnected	src-gmu/Grenade/OpenCL/FullyConnected.hs:32:5-9	857	3213	0.0	0.0	0.0	0.0
writeBufferR	Grenade.OpenCL.Context	src-gmu/Grenade/OpenCL/Context.hs:(124,1)-(129,46)	850	3213	0.1	0.0	0.1	0.0
clEnqueueWriteBuffer'	Grenade.OpenCL.Context	src-gmu/Grenade/OpenCL/Context.hs:(223,1)-(226,41)	858	3213	0.0	0.0	0.0	0.0
clQueue	Grenade.OpenCL.Context	src-gmu/Grenade/OpenCL/Context.hs:59:5-11	856	3213	0.0	0.0	0.0	0.0
unId	Data.Vector.Fusion.Util	Data/Vector/Fusion/Util.hs:25:21-24	855	3213	0.0	0.0	0.0	0.0
netTrain	Main	main/feedforward.hs:(38,1)-(53,53)	802	1	14.8	8.9	99.2	98.9
unId	Data.Vector.Fusion.Util	Data/Vector/Fusion/Util.hs:25:21-24	817	1300000	0.0	0.0	0.0	0.0
train	Grenade.Core.Runner	src/Grenade/Core/Runner.hs:(50,1)-(52,38)	812	100000	0.3	0.0	81.0	79.8
applyUpdate	Grenade.Core.Network	src/Grenade/Core/Network.hs:(152,1)-(156,8)	840	700000	8.8	6.3	15.3	11.3
lpRef	Grenade.OpenCL.FullyConnected	src-gmu/Grenade/OpenCL/FullyConnected.hs:31:5-9	841	300003	0.0	0.0	0.0	0.0
clEnqueueNDRangeKernel'	Grenade.OpenCL.Context	src-gmu/Grenade/OpenCL/Context.hs:(229,1)-(236,35)	848	300000	6.0	4.6	6.0	4.6
clQueue	Grenade.OpenCL.Context	src-gmu/Grenade/OpenCL/Context.hs:59:5-11	849	300000	0.5	0.4	0.5	0.4
kUpdate	Grenade.OpenCL.FullyConnected	src-gmu/Grenade/OpenCL/FullyConnected.hs:28:5-11	845	300000	0.0	0.0	0.0	0.0
clSetKernelArgSto'	Grenade.OpenCL.Context	src-gmu/Grenade/OpenCL/Context.hs:(239,1)-(240,79)	844	300000	9	0.0	0.0	0.0
learningMomentum	Grenade.Core.LearningParameters	src/Grenade/Core/LearningParameters.hs:19:5-20	846	3	0.0	0.0	0.0	0.0
learningRate	Grenade.Core.LearningParameters	src/Grenade/Core/LearningParameters.hs:18:5-16	842	3	0.0	0.0	0.0	0.0
learningRegulariser	Grenade.Core.LearningParameters	src/Grenade/Core/LearningParameters.hs:20:5-23	847	3	0.0	0.0	0.0	0.0
backPropagate	Grenade.Core.Runner	src/Grenade/Core/Runner.hs:(37,1)-(40,13)	813	100000	1.0	0.3	65.4	68.5
runGradient	Grenade.Core.Network	src/Grenade/Core/Network.hs:(131,1)-(144,17)	815	100000	4.1	12.2	4.1	12.2
runNetwork	Grenade.Core.Network	src/Grenade/Core/Network.hs:(103,1)-(116,17)	814	0	30.2	31.7	60.4	56.1
wTape	Grenade.OpenCL.FullyConnected	src-gmu/Grenade/OpenCL/FullyConnected.hs:32:5-9	825	600000	0.5	0.8	0.5	0.8
unId	Data.Vector.Fusion.Util	Data/Vector/Fusion/Util.hs:25:21-24	839	500000	0.0	0.0	0.0	0.0
clEnqueueNDRangeKernel'	Grenade.OpenCL.Context	src-gmu/Grenade/OpenCL/Context.hs:(229,1)-(236,35)	827	300000	5.7	4.6	5.7	4.6
clQueue	Grenade.OpenCL.Context	src-gmu/Grenade/OpenCL/Context.hs:59:5-11	828	300000	0.3	0.4	0.3	0.4
kForward	Grenade.OpenCL.FullyConnected	src-gmu/Grenade/OpenCL/FullyConnected.hs:29:5-12	829	300000	0.3	0.4	0.3	0.4
readBufferR	Grenade.OpenCL.Context	src-gmu/Grenade/OpenCL/Context.hs:(132,1)-(141,46)	830	300000	10.8	12.1	15.1	14.8
unId	Data.Vector.Fusion.Util	Data/Vector/Fusion/Util.hs:25:21-24	835	1100000	0.0	0.0	0.0	0.0
clEnqueueReadBuffer'	Grenade.OpenCL.Context	src-gmu/Grenade/OpenCL/Context.hs:(216,1)-(219,41)	838	300000	4.4	2.7	4.4	2.7
clQueue	Grenade.OpenCL.Context	src-gmu/Grenade/OpenCL/Context.hs:59:5-11	836	300000	0.0	0.0	0.0	0.0
wOut	Grenade.OpenCL.FullyConnected	src-gmu/Grenade/OpenCL/FullyConnected.hs:34:5-8	837	300000	0.2	0.4	0.2	0.4
writeBufferR	Grenade.OpenCL.Context	src-gmu/Grenade/OpenCL/Context.hs:(124,1)-(129,46)	816	300000	2.8	0.4	8.1	3.1
clEnqueueWriteBuffer'	Grenade.OpenCL.Context	src-gmu/Grenade/OpenCL/Context.hs:(223,1)-(226,41)	826	300000	5.3	2.7	5.3	2.7
clQueue	Grenade.OpenCL.Context	src-gmu/Grenade/OpenCL/Context.hs:59:5-11	824	300000	0.0	0.0	0.0	0.0
unId	Data.Vector.Fusion.Util	Data/Vector/Fusion/Util.hs:25:21-24	823	300000	0.0	0.0	0.0	0.0
getStdRandom	System.Random	System/Random.hs:(586,1)-(587,26)	803	0	3.3	10.2	3.3	10.2
randomNet	Main	main/feedforward.hs:35:1-25	705	1	0.0	0.0	0.0	0.0
randomFullyConnectedCL	Grenade.OpenCL.FullyConnected	src-gmu/Grenade/OpenCL/FullyConnected.hs:(99,1)-(145,31)	706	3	0.0	0.0	0.0	0.0
unId	Data.Vector.Fusion.Util	Data/Vector/Fusion/Util.hs:25:21-24	774	592	0.0	0.0	0.0	0.0
clSetKernelArgSto'	Grenade.OpenCL.Context	src-gmu/Grenade/OpenCL/Context.hs:(239,1)-(240,79)	801	63	0.0	0.0	0.0	0.0
mkBufferR	Grenade.OpenCL.Context	src-gmu/Grenade/OpenCL/Context.hs:(158,1)-(163,46)	769	50	0.0	0.0	0.0	0.0
unId	Data.Vector.Fusion.Util	Data/Vector/Fusion/Util.hs:25:21-24	780	50	0.0	0.0	0.0	0.0
clContext	Grenade.OpenCL.Context	src-gmu/Grenade/OpenCL/Context.hs:57:5-13	772	18	0.0	0.0	0.0	0.0


```

instance (KnownNat i, KnownNat o) => UpdateLayer (FullyConnectedCL i o) where
  type Gradient (FullyConnectedCL i o) = (FullyConnected'CL i o)

  -- let (nB, nBM) = descendVector lp oB bG oBM
  -- (nA, nM) = descendMatrix lp oA aG oM
  -- in FullyConnectedCL (FullyConnected'CL nB nA) (FullyConnected'CL nBM nM)
runUpdate lp n _ =
  unsafeWithCL $ \cl -> do
    let k = kUpdate n
    lastLp <- readIORef (lpRef n)
    unless (Just lp == lastLp) $ do
      clSetKernelArgSto' k 2 (learningRate lp)
      clSetKernelArgSto' k 3 (learningMomentum lp)
      clSetKernelArgSto' k 4 (learningRegulariser lp)
      writeIORef (lpRef n) (Just lp)
    _ <- clEnqueueNDRangeKernel' (clQueue cl) k [len]
    pure n
  where
    len = fromIntegral (natVal (Proxy @o))

createRandom = randomFullyConnectedCL

instance (KnownNat i, KnownNat o) => Layer (FullyConnectedCL i o) ('D1 i) ('D1 o) where
  type Tape (FullyConnectedCL i o) ('D1 i) ('D1 o) = CLBuffer (R i)

  -- Do a matrix vector multiplication and return the result.
  -- (v, S1D (wB + wN #> v))
runForwards n (S1D v) =
  unsafeWithCL $ \cl -> do
    writeBufferR cl (wTape n) v
    _ <- clEnqueueNDRangeKernel' (clQueue cl) (kForward n) [len]
    res <- readBufferR cl (wOut n) (Just (konst 0))
    pure (wTape n, S1D res)
  where
    len = fromIntegral (natVal (Proxy @o))

  -- Run a backpropogation step for a full connected layer.
  -- let wB' = dEdy; mm' = dEdy `outer` x
  -- dWs = tr wN #> dEdy
  -- in (FullyConnected'CL wB' mm', S1D dWs)
runBackwards n _ (S1D dEdy) =
  unsafeWithCL $ \cl -> do
    let FullyConnected'CL bG _ = wGradient n
    writeBufferR cl bG dEdy
    _ <- clEnqueueNDRangeKernel' (clQueue cl) (kBackward n) [len]
    dWs <- readBufferR cl (wIn n) (Just (konst 0))
    pure (wGradient n, S1D dWs)
  where
    len = fromIntegral (max (natVal (Proxy @o)) (natVal (Proxy @i)))

```

```

{
  inputs.nixpkgs.url = github:nixos/nixpkgs/master;
  inputs.CLUtil = { url = github:acowley/CLUtil/master; flake = false; };

  outputs = { self, nixpkgs, CLUtil }: let
    inherit (nixpkgs.lib) flip mapAttrs mapAttrsToList;
    inherit (pkgs.nix-gitignore) gitignoreSourcePure;

    pkgs = import nixpkgs {
      system = "x86_64-linux";
      config.allowUnfree = true;
      overlays = [ self.overlay ];
    };
    hsPkgs = pkgs.haskellPackages;
    getSrc = dir: gitignoreSourcePure [./.gitignore] dir;
  in {
    overlay = final: prev: let
      inherit (prev.haskell.lib) doJailbreak dontCheck unmarkBroken overrideCabal dontHaddock;
    in {
      haskell = prev.haskell // {
        packageOverrides = prev.lib.composeExtensions
          (prev.haskell.packageOverrides or {}: {}))
          (hself: hsuper: {
            grenade = hself.callCabal2nix "grenade" (getSrc ./grenade) {};
            CLUtil = dontHaddock (dontCheck (hself.callCabal2nix "CLUtil" CLUtil {}));
            OpenCL = overrideCabal (dontCheck (doJailbreak (unmarkBroken hsuper.OpenCL))) (drv: {
              configureFlags = (drv.configureFlags or []) ++ [
                "--extra-lib-dirs=${pkgs.ocl-icd}/lib"
                "--extra-include-dirs=${pkgs.opencl-headers}/include"
              ];
            });
            profiteur = doJailbreak (unmarkBroken hsuper.profitreur);
          });
      };
    };

    devShell.x86_64-linux = hsPkgs.shellFor {
      withHoogle = true;
      packages = p: [ p.grenade ];
      buildInputs = [
        hsPkgs.cabal-install
        hsPkgs.haskell-language-server
        hsPkgs.stylish-haskell
        hsPkgs.profitreur
      ];
    };
  };
}

```